



(A) Mixed audio — both speakers overlap at $r=0.05$
Whisper transcribes correctly: CER = 0.44



(B) Oracle-separated Speaker 1 (con_006)
Speech ends ~2.1s, then SILENCE — Whisper transcribes OK: CER = 0.44



(C) Oracle-separated Speaker 2 (pro_006) — CATASTROPHIC HALLUCINATION
Leading silence → Whisper enters token-id repetition loop: CER = 24.25, CR = 16.33 (transcript 24× longer than reference!)

